

Machine Learning Nanodegree Capstone Project

Kaggle Challenge: PetFinder.my Adoption Prediction

Mahesh Babu Gorantla - March 15, 2019

- **Machine Learning Nanodegree Capstone Project**
 - **Kaggle Challenge: PetFinder.my Adoption Prediction**
 - **Project Overview**
 - **Problem Statement**
 - **Metrics**
 - **Data Exploration**
 - **Data Fields**
 - **AdoptionSpeed**
 - **Images**
 - **Image Metadata**
 - **Sentiment Data**
 - **Dropping Irrelevant Columns**
 - **Missing Values Per Column**
 - **Exploratory Visualization**
 - **Descriptive Statistics**
 - **Dataset Distribution (per class)**
 - **Algorithms and Techniques**
 - **Benchmark**
 - **Data Preprocessing**
 - **Implementation**
 - **Model Evaluation and Validation**
 - **Justification**
 - **Free-Form Visualization**
 - **Reflection**
 - **Functionality**
 - **References**

Project Overview

Student provides a high-level overview of the project in layman's terms. Background information such as the problem domain, the project origin, and related data sets or input data is given.

Each year, approximately 7.6 million companion animals enter the animal shelters nationwide (ASPCA). Of those, approximately 3.9 million are dogs and 3.4 million are cats. About 2.7 million shelter animals are adopted each year (1.4 million dogs and 1.3 million cats). This leaves around two-thirds of pets going unadopted.

This uncertainty of the unadopted pets future is further worsened because many a times a pet shelter's future is typically uncertain and many pet shelters cannot keep the unadopted pets for extended periods of time due to logistical reasons and many of the pet shelters are run voluntarily (mostly dependent on generous donations).

If at all we can find those stray animals a forever home, many precious lives can be given a new life that they truly deserve and in turn we will see more happier families.

I am big fan of pets. My personal motivation to solve this problem explicitly is because doing so will make the [PetFinder.my](#) Agency streamline their pet adoption services which in turn help more pets being adopted. Being able to predict how quickly a pet will be adopted will help pet adoption agencies optimize their logistics and operate more efficiently reducing the percentage of unadopted pets.

Problem Statement

The problem which needs to be solved is clearly defined. A strategy for solving the problem, including discussion of the expected solution, has been made.

[PetFinder.my](#) has been Malaysia's leading animal welfare platform since 2008, with a database of more than 150,000 animals. PetFinder collaborates closely with animal lovers, media, corporations, and global organizations to improve animal welfare.

The problem at hand currently is to create a model to predict the adoption speed of a pet given its features (mentioned below).

Using a machine learning algorithm(s) to predict the adoption speed of a pet should help us solve the above mentioned problem.

The Pet Adoption Prediction is a multi-class classification problem. In this problem, I need to develop algorithms to predict the adoptability of pets - specifically, how quickly is a pet adopted based on its features (which have been explained in the section below).

It is a classification problem because a pet's adoption speed is classified into 5 categories.

A relevant potential solution is to use a few classification algorithms such as

- k-Means Clustering
- DBSCAN
- Random Forests

- XGBoost (Baseline model)

And, choose the one with best `cross_validation` score on the training dataset.

Relevant Academic Research has been cited in the `References` section

Metrics

Metrics used to measure performance of a model or result are clearly defined. Metrics are justified based on the characteristics of the problem.

I would like to propose an evaluation metric called `Cross Validation Score`. The simplest way to use cross-validation is to call the `cross_val_score` helper function on the estimator and the dataset.

Given, the problem is to develop a supervised multi-class classification model, the `cross_val_score` method uses `StratifiedKFold` cross-validation technique. `StratifiedKFold` cross-validation method is a variation of `KFold` that returns stratified folds. The folds are made by preserving the percentage of samples for each class.

The model whose cross validation score is greater than the benchmark model will be my solution model.

Data Exploration

If a dataset is present, features and calculated statistics relevant to the problem have been reported and discussed, along with a sampling of the data. In lieu of a dataset, a thorough description of the input space or input data has been made. Abnormalities or characteristics about the data or input that need to be addressed have been identified.

The dataset to solve the above mentioned problem has been obtained as a part of Kaggle Challenge. In this problem, I will predict the speed at which a pet is adopted, based on the pet's listing on PetFinder. Sometimes a profile represents a group of pets. In this case, the speed of adoption is determined by the speed at which all of the pets are adopted. The data included text, tabular, and image data.

Data Fields

- `PetID` - Unique hash ID of pet profile
- `AdoptionSpeed` - Categorical speed of adoption. Lower is faster. `This is the value to predict`. See below section for more info.
- `Type` - Type of animal (1 = Dog, 2 = Cat)
- `Name` - Name of pet (Empty if not named)
- `Age` - Age of pet when listed, in months
- `Breed1` - Primary breed of pet (Refer to BreedLabels dictionary)
- `Breed2` - Secondary breed of pet, if pet is of mixed breed (Refer to BreedLabels dictionary)
- `Gender` - Gender of pet (1 = Male, 2 = Female, 3 = Mixed, if profile represents group of pets)
- `Color1` - Color 1 of pet (Refer to ColorLabels dictionary)
- `Color2` - Color 2 of pet (Refer to ColorLabels dictionary)
- `Color3` - Color 3 of pet (Refer to ColorLabels dictionary)
- `MaturitySize` - Size at maturity (1 = Small, 2 = Medium, 3 = Large, 4 = Extra Large, 0 = Not Specified)
- `FurLength` - Fur length (1 = Short, 2 = Medium, 3 = Long, 0 = Not Specified)
- `Vaccinated` - Pet has been vaccinated (1 = Yes, 2 = No, 3 = Not Sure)
- `Dewormed` - Pet has been dewormed (1 = Yes, 2 = No, 3 = Not Sure)
- `Sterilized` - Pet has been spayed / neutered (1 = Yes, 2 = No, 3 = Not Sure)
- `Health` - Health Condition (1 = Healthy, 2 = Minor Injury, 3 = Serious Injury, 0 = Not Specified)
- `Quantity` - Number of pets represented in profile Fee - Adoption fee (0 = Free)
- `State` - State location in Malaysia (Refer to StateLabels dictionary)
- `RescuerID` - Unique hash ID of rescuer
- `VideoAmt` - Total uploaded videos for this pet
- `PhotoAmt` - Total uploaded photos for this pet
- `Description` - Profile write-up for this pet. The primary language used is English, with some in Malay or Chinese.

AdoptionSpeed

Contestants are required to predict this value. The value is determined by how quickly, if at all, a pet is adopted. The values are determined in the following way:

- 0 - Pet was adopted on the same day as it was listed.
- 1 - Pet was adopted between 1 and 7 days (1st week) after being listed.
- 2 - Pet was adopted between 8 and 30 days (1st month) after being listed.
- 3 - Pet was adopted between 31 and 90 days (2nd & 3rd month) after being listed.
- 4 - No adoption after 100 days of being listed. (There are no pets in this dataset that waited between 90 and 100 days)

Images

For pets that have photos, they will be named in the format of *PetID-ImageNumber.jpg*. Image \$1\$ is the profile (default) photo set for the pet. For privacy purposes, faces, phone numbers and emails have been masked.

Image Metadata

We have run the images through **Google's Vision API**, providing analysis on **Face Annotation**, **Label Annotation**, **Text Annotation** and **Image Properties**. You may optionally utilize this supplementary information for your image analysis.

File name format is *PetID-ImageNumber.json*.

Some properties will not exist in JSON file if not present, i.e. Face Annotation. Text Annotation has been simplified to just 1 entry of the entire text description (instead of the detailed JSON result broken down by individual characters and words). Phone numbers and emails are already anonymized in Text Annotation.

Google Vision API reference: <https://cloud.google.com/vision/docs/reference/rest/v1/images/annotate>

Sentiment Data

We have run each pet profile's description through **Google's Natural Language API**, providing analysis on sentiment and key entities. You may optionally utilize this supplementary information for your pet description analysis. There are some descriptions that the API could not analyze. As such, there are fewer sentiment files than there are rows in the dataset.

File name format is *PetID.json*.

Google Natural Language API reference: <https://cloud.google.com/natural-language/docs/basics>

Dropping Irrelevant Columns

- **Name, PetID, RescuerID**
 - It makes sense to not have **Name** column because it does not help us predict how faster a pet animal will be adopted as we know that any prospective pet owner will not adopt a pet animal based on its name.
 - Likewise **PetID, RescuerID** also have no relevance to the task at hand
- However, for now to simplify the modeling let's ignore the following columns
 - **Fee** - Might have some implication if a pet will adopted (Need to check the correlation)
 - **State** - Might not influence the adoption speed (Need to check the correlation)
 - **VideoAmt, PhotoAmt & Description** - Will have some impact on the adoption speed. Because a beautiful looking pet might get adopted sooner. But just to make the modeling simpler let's ignore these columns as well for now.

Missing Values Per Column

A common abnormality found in datasets is **missing values**

Column Name	No. of Missing Values
Type	0
Name	1257
Age	0
Breed1	0
Breed2	0
Gender	0
Color1	0
Color2	0
Color3	0
FurLength	0
Vaccinated	0
Dewormed	0
Sterilized	0

Column Name	No. of Missing Values
Health	0
Quantity	0
Fee	0
State	0
RescuerID	0
VideoAmt	0
Description	12
PetID	0
PhotoAmt	0
AdoptionSpeed	0

Very often rows with missing values will be dropped as anomalies in the data however, given that the dataset only has 14,993 datapoints removing 1269 rows we are almost throwing away nearly 10% of the data which will not be an ideal operation to perform as a part of data cleaning process.

However as mentioned above, the columns Name and Description are being dropped as irrelevant columns for the problem at hand. Hence, we are effectively removing anomalies instead of completely deleting the other columns that are needed.

Exploratory Visualization

A visualization has been provided that summarizes or extracts a relevant characteristic or feature about the dataset or input data with thorough discussion. Visual cues are clearly defined.

Descriptive Statistics

Glimpse of Descriptive Statistics of the dataset

	Type	Age	Breed1	Breed2	Gender	Color1	Color2	Color3	MaturitySize	FurLength	Vaccinated	Dewormed	Ste
count	14993	14993	14993	14993	14993	14993	14993	14993	14993	14993	14993	14993	14993
mean	1.45761	0.871006	265.273	74.0097	1.77616	2.23418	3.22284	1.88201	1.862	1.46748	1.73121	1.55873	1.9
std	0.498217	1.51298	60.0568	123.012	0.681592	1.74523	2.74256	2.98409	0.547959	0.59907	0.667649	0.695817	0.5
min	1	0	0	0	1	1	0	0	1	1	1	1	1
25%	1	0.166667	265	0	1	1	0	0	2	1	1	1	2
50%	1	0.25	266	0	2	2	2	0	2	1	2	1	2
75%	2	1	307	179	2	3	6	5	2	2	2	2	2
max	2	21.25	307	307	3	7	7	7	4	3	3	3	3

Dataset Distribution (per class)

Adoption Speed	Data Points Size
0	410
1	3090
2	4037
3	3259
4	4197

From the above observations we can determine that the dataset is almost imbalanced.

Algorithms and Techniques

Algorithms and techniques used in the project are thoroughly discussed and properly justified based on the characteristics of the problem.

Benchmark

Student clearly defines a benchmark result or threshold for comparing performances of solutions obtained.

Data Preprocessing

All preprocessing steps have been clearly documented. Abnormalities or characteristics about the data or input that needed to be addressed have been corrected. If no data preprocessing is necessary, it has been clearly justified.

Implementation

The process for which metrics, algorithms, and techniques were implemented with the given datasets or input data has been thoroughly documented. Complications that occurred during the coding process are discussed.

Model Evaluation and Validation

The final model's qualities — such as parameters — are evaluated in detail. Some type of analysis is used to validate the robustness of the model's solution.

Justification

The final results are compared to the benchmark result or threshold with some type of statistical analysis. Justification is made as to whether the final model and solution is significant enough to have adequately solved the problem.

Free-Form Visualization

A visualization has been provided that emphasizes an important quality about the project with thorough discussion. Visual cues are clearly defined.

Reflection

Student adequately summarizes the end-to-end problem solution and discusses one or two particular aspects of the project they found interesting or difficult.

Functionality

Discussion is made as to how one aspect of the implementation could be improved. Potential solutions resulting from these improvements are considered and compared/contrasted to the current solution.

References

1. Weiss, E., Miller, K., Mohan-Gibbons, H., & Vela, C. (2012). Why Did You Choose This Pet?: Adopters and Pet Selection Preferences in Five Animal Shelters in the United States. *Animals*, 2(2), 144-159. doi:10.3390/ani2020144
2. Weiss, E., Dolan, E., Garrison, L., Hong, J., & Slater, M. (2013). Should Dogs and Cats be Given as Gifts? *Animals*, 3(4), 995-1001. doi:10.3390/ani3040995
3. Garrison, L., & Weiss, E. (2014). What Do People Want? Factors People Consider When Acquiring Dogs, the Complexity of the Choices They Make, and Implications for Nonhuman Animal Relocation Programs. *Journal of Applied Animal Welfare Science*, 18(1), 57-73. doi:10.1080/10888705.2014.943836
4. PetFinder.my Adoption Prediction. (n.d.). Retrieved from <https://www.kaggle.com/c/petfinder-adoption-prediction>
5. Raschka, S. (2014, August 03). Linear Discriminant Analysis. Retrieved from https://sebastianraschka.com/Articles/2014_python_lda.html
6. Cross-validation: Evaluating estimator performance¶. (n.d.). Retrieved from https://scikit-learn.org/stable/modules/cross_validation.html
7. Sklearn.metrics.f1_score¶. (n.d.). Retrieved from https://scikit-learn.org/stable/modules/generated/sklearn.metrics.f1_score.html#sklearn.metrics.f1_score.html
8. Sklearn.model_selection.cross_val_score¶. (n.d.). Retrieved from https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.cross_val_score.html
9. Sklearn.model_selection.StratifiedKFold¶. (n.d.). Retrieved from https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.StratifiedKFold.html#sklearn.model_selection.StratifiedKFold.html
10. Sklearn.base.TransformerMixin¶. (n.d.). Retrieved from <https://scikit-learn.org/stable/modules/generated/sklearn.base.TransformerMixin.html>
11. Ozdemir, S., & Susarla, D. (2018). Feature engineering made easy: Identify unique features from your dataset in order to build powerful machine learning systems. Birmingham, UK: Packt Publishing